

John Piontkowsky  
Vahe Kilamyan  
April Luo  
Marc Mendiola  
Gueho Choi

## TFS Part 2

### Design Decisions

#### 1. Master

- a. Multi-threaded
- b. Spawns a new thread to service each new request
- c. Contains:
  - i. FileSystem.java class (creates, removes, and interacts with directories and files)
  - ii. Hashtable to map Strings to a TFSDirectory object
  - iii. TFSDirectory class (contains a list of subdirectories and the files within those subdirectories)
    - 1. ex. “\” represents the root directory
    - 2. Stored in a sorted set; lookup time =  $O(\log n)$
- d. Each transaction is logged using FSLogger.java, for crash recovery purposes
- e. All transactions other than Read, Write, and Append handled by Master itself
- f. Global unique identifiers: Master hashes the full file path and passes the resulting hash to the chunkservers
- g. ChunkTracker.java (within Master.java)
  - i. Maintains a list of live chunkservers and a list of dead chunkservers
  - ii. Sends out a “heartbeat” every 60 seconds to ensure existing chunkserver connections are still alive
  - iii. If no response, moves the chunkserver to the dead list

#### 2. Chunkserver

- a. Each chunkserver “registers” with Master on startup
- b. Registering with Master gives the chunkserver a number (eg, 1,2).
- c. Chunkserver’s files are stored in the folder cs + its number (eg, “cs1”)
- d. Chunkserver attempts to synchronize with other chunkservers by examining its folder for .meta files, which contain the version number of files and the other chunkservers containing them.
- e. Updates its files by issuing “updateFile” commands to other chunkservers contained in the .meta file.

+

### 3. Client

- a. Single-threaded
- b. Forwards requests to either Master or chunkserver

#### Interfaces and Filesystem Calls

	Input	Output	Data flow
<b>CreateFile</b>	(String fileName, String initialData)	File created on chunkserver using Java File Create and data passed in	<p>Client to Master, Master to Client, Client to Chunkserver</p> <p><u>Client to Master</u>: Client sends create request to master, specifying a unique file name.</p> <p><u>Master to Client</u>: Master ensures that filename is unique in this directory. If it is, Master chooses n chunkservers and returns their addresses to the client. If it is not Master returns an error to the client.</p> <p><u>Client to Chunkserver</u>: Client chooses 1 of n chunkservers and issues a create request to the chunkserver. Chunkserver uses fileName and initialData to create the file. Client proceeds to repeat this process n times until the specified number of replicas have been created.</p>
<b>CreateDir</b>	(String dirName)	Directory created on chunkserver. Return -1 on failure.	<p>Client to Master, Master to Client</p> <p><u>Client to Master</u>: Client sends create directory request to master, specifying a unique directory name.</p> <p><u>Master to Client</u>: Master ensures that directory is unique. If it is, it creates the directory and returns success to the client. Otherwise, it returns failure.</p>

<b>Open</b>	(String filename)	Returns List of chunkservers holding a replica of this file.	<p>Client to Master, Master to Client, Client to Chunkserver</p> <p><u>Client to Master</u>: Client sends an open request to master, specifying a unique file name.</p> <p><u>Master to Client</u>: Master ensures that the file is valid (exists, is unique in directory). If it is, Master sends back a structure that contains the IP addresses of chunk servers with this file on it on and the chunk value.</p>
<b>Close</b>	(String filename)	Returns boolean indicating success	<p>Client to Master</p> <p><u>Client to Master</u>: Client sends a close request to master.</p> <p><u>Master to Client</u>: Master sends back a confirmation indicating success.</p>
<b>Read</b>	(String filename, int filePosition, int numBytes)	Returns struct of data and an error code indicating success or failure.	<p>Client to Chunkserver, Chunkserver to Client</p> <p><u>Client to Chunkserver</u>: Client sends a read request to chunkserver, specifying a chunk handle.</p> <p><u>Chunkserver to Client</u>: Chunkserver attempts to read the chunk at the chunk handle. If the chunk exists, the chunkserver sends the data back to the client.</p>
<b>Write</b>	(String filename, int offset, File initialData)	Boolean representing success or failure of write operation.	<p>Client to Chunkserver, Chunkserver to Client</p> <p><u>Client to Chunkserver</u>: Client uses the chunkserver selected through the open operation as an address of</p>

			<p>the request. Client then sends message to chunkserver specifying filename, file offset location, and data to be written. Chunkserver will then perform the write operation on the file specified.</p> <p><u>Chunkserver to Client:</u> After successful execution of the write operation chunkserver will send a boolean success(true) to client. Otherwise the chunkserver will send a boolean failure(false).</p>
<b>DeleteFile</b>	(String filename)	Returns boolean success of the delete operation, true if success, false otherwise.	<p>Client to Master, Master to Client, Client to Chunkserver, Chunkserver to Client</p> <p><u>Client to Master:</u> Client sends filename of file to be deleted to Master.</p> <p><u>Master to Client:</u> Master returns a list of chunkservers which contain the file to be deleted.</p> <p><u>Client to Chunkserver:</u> Client sends each member of the list a message to delete the file with the provided filename.</p> <p><u>Chunkserver to Client:</u> Chunkserver sends confirmation of completion to client after the operation has been completed.</p>
<b>FileSize</b>	(String filename)	Returns integer representation of the file specified	<p>Client to Master, Master to Client, Client to Chunkserver, Chunkserver to Client</p> <p><u>Client to Master:</u> Client sends the filename to master to determine</p>

			<p>location of the file on chunkservers.</p> <p><u>Master to Client</u>: Master returns a List of chunkservers to the client which contain the file specified.</p> <p><u>Client to Chunkserver</u>: Client chooses one of the specified chunkservers to query the size of the file on that chunkserver.</p>
<b>LS</b>	(String directoryname)	Returns data representation of the current file structure being tracked by master.	<p>Client to Master, Master to Client</p> <p><u>Client to Master</u>: Client sends the request to master specifying a directory name which it would like to receive the contents of.</p> <p><u>Master to Client</u>: Master takes the directory name and compares it to its internal file structure records. Master then returns the subdirectories and appropriate files contained within the directory specified.</p>